



Lecture 2

Database System Concepts and Architecture

Agenda



Data Models



Categories of Data Models



Schema



Database state



Three-Schema Architecture



DBMS Languages



Client/Server Architectures for DBMSs



Data Models

Data Models

- One fundamental characteristic of the database approach is that it provides some level of data abstraction.
- **Data abstraction** generally refers to:
 - **Suppression of details** of data organization and storage.
 - **Highlighting of the essential features** for an improved understanding of data.

Data Models

- **Data abstraction** is achieved by a **Data model**.
- What is a Data Model?
 - A collection of **concepts** that can be used to describe the *structure of a database*.
 - By *structure of a database* we mean the **data types, relationships, and constraints** that apply to the data.
 - Most data models also include a set of **basic operations** for specifying *retrievals* and *modifications* on the database.



Categories of Data Models

Categories of Data Models

Conceptual data
models
(High Level)

Representational
data models
(Implementational)

Physical data
models
(Low Level)

Categories of Data Models

Conceptual data
models
(High Level)



Provide concepts that
are close to the way
many users perceive
data

Representational
data models
(Implementational)



Provide concepts that
may be easily
understood by end users
and **aren't too far** from
how data is organized in
computer storage

Physical data
models
(Low Level)



Provide concepts that
describe the details of
how data is stored on
the computer storage
media

Categories of Data Models

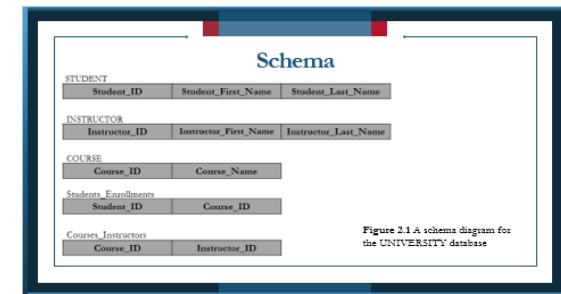
- Conceptual data models → **Entity–Relationship model**
- Representational data models → **Relational data model**
- Physical data models → **Access path**



Schema

Schema

- In a data model, it is important to distinguish between the *description* of the database and the *database itself*.
- The description of a database is called the **database schema**, which is specified during database design and is not expected to change frequently.
- A schema diagram for the UNIVERSITY database given before:



Schema

STUDENT

Student_ID	Student_First_Name	Student_Last_Name
------------	--------------------	-------------------

INSTRUCTOR

Instructor_ID	Instructor_First_Name	Instructor_Last_Name
---------------	-----------------------	----------------------

COURSE

Course_ID	Course_Name
-----------	-------------

Students_Enrollments

Student_ID	Course_ID
------------	-----------

Courses_Instructors

Course_ID	Instructor_ID
-----------	---------------

Figure 2.1 A schema diagram for the UNIVERSITY database

Schema

- The schema diagram displays the **structure of each record** but not the **actual records**.
- A schema diagram displays only *some aspects* of a schema, such as the names of record types and data items, and *some types of constraints*.
- Other aspects are not specified in the schema diagram; for example, the **data type** of each data item isn't present.



Database state

Database state

- The actual data in a database may change frequently.
- For example, the UNIVERSITY database changes every time we add a new student or add a new course.
- The data in the database at a particular moment in time is called a **database state** or **snapshot**.

Database state

- Every time we insert or delete a record or change the value of a data item in a record, we **change one state of the database into another state**.
- When we **define** a new database, we specify its database schema only to the DBMS.
- At this point, the corresponding database state is the *empty state* with no data.

Database state

- We get the *initial state* of the database when the database is first populated or loaded with the initial data.
- From then on, every time a modification operation is applied to the database, we get another database state.
- At any point in time, the database has a *current state*.

Database state

- The DBMS is partly responsible for ensuring that every state of the database is a **valid state**—that is, a state that satisfies the structure and constraints specified in the schema.
- The DBMS stores the descriptions of the schema constructs and constraints—also called the **meta-data**—in the DBMS catalog so that DBMS software can refer to the schema whenever it needs to.



Three-Schema Architecture

Three-Schema Architecture

- The goal of the three-schema architecture is to separate the user applications from the physical database.
- In this architecture, schemas can be defined at the following three levels:
 - **Internal Level**
 - **Conceptual Level**
 - **External or View Level**

Three-Schema Architecture

➤ Internal Level

- Has an **internal schema**, which describes the physical storage structure of the database.
- The internal schema uses a **physical data model** and describes the complete details of data storage and access paths for the database.

Three-Schema Architecture

➤ Conceptual Level

- Has a **conceptual schema**, which describes the structure of the whole database for users.
- The conceptual schema hides the details of physical storage structures and concentrates on describing entities, data types, relationships, user operations, and constraints.
- Usually, a **representational data model** is used to describe the conceptual schema.

Three-Schema Architecture

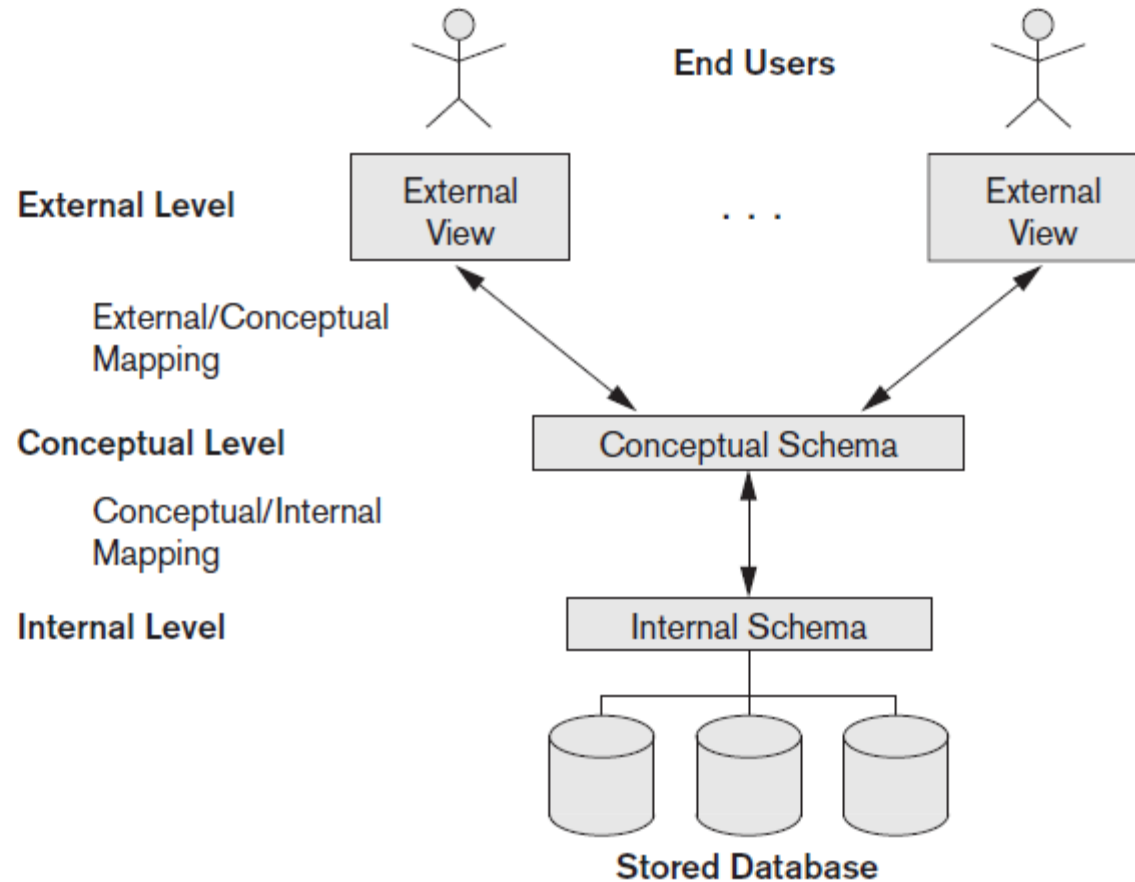
➤ External or View Level

- Includes a number of **external schemas** or **user views**.
- Each external schema describes the part of the database that a particular user group is interested in and hides the rest of the database from that user group.
- Each external schema is typically uses a **representational data model**.

Three-Schema Architecture

Figure 2.2

The three-schema architecture.



Three-Schema Architecture

- For the UNIVERSITY database given, show an example of a **Conceptual Schema** and an **External Schema**.



DBMS Languages

DBMS Languages

- In many DBMSs where no strict separation of levels is maintained, one language, called the **data definition language (DDL)**, is used to define conceptual and internal schemas.
- In DBMSs where a clear separation is maintained between the conceptual and internal levels, the DDL is used to specify the conceptual schema only. Another language, the **storage definition language (SDL)**, is used to specify the internal schema.

DBMS Languages

- In most relational DBMSs today, there *is no specific language* that performs the role of SDL. Instead, the internal schema is specified by a combination of functions, parameters, and specifications related to storage of files.
- For a true three-schema architecture, we would need a third language, the **view definition language (VDL)** to specify user views and their mappings to the conceptual schema.

DBMS Languages

- In most DBMSs *the DDL is used to define both conceptual and external schemas.*
- In relational DBMSs, **Structured Query Language (SQL)** is used in the role of VDL to define user or application views.
- Once the database schemas are compiled and the database is populated with data, users must have some means to manipulate the database (i.e. *retrieval, insertion, deletion, and updating* of the data). DBMS provides a language called the **data manipulation language (DML)** for these purposes.

DBMS Languages

- **DDL:** Conceptual schemas
Internal schemas (if no separation)
External schemas
- **SDL:** Internal schemas
- **VDL:** User views and their mappings
- **SQL:** User or application views
- **DML:** Manipulate the database

DBMS Languages

- In current DBMSs, the preceding types of languages are usually *not considered distinct languages*; rather, a comprehensive integrated language is used that includes constructs for conceptual schema definition, view definition, and data manipulation. Storage definition is typically kept separate.
- A typical example of a comprehensive database language is the **SQL** relational database language, which represents a combination of *DDL*, *VDL*, and *DML*, as well as statements for constraint specification, schema evolution, and many other features.



Client/Server Architectures for DBMSs

Client/Server Architectures for DBMSs

- DBMS was a **centralized DBMS** in which all the DBMS functionality, application program execution, and user interface processing were carried out on one machine.
- Gradually, DBMS systems started to exploit the available processing power at the user side, which led to **client/server DBMS** architectures.

Client/Server Architectures for DBMSs

- The idea is to define **specialized servers** with specific functionalities.
- For example, we can have:
 - **File server** that maintains the files.
 - **Printer server** connected to various printers; all print requests by the clients are forwarded to this machine.
 - **Web servers** or **e-mail servers**.
- The **client machines** provide the user with the appropriate interfaces to utilize these servers as well as with local processing power to run local applications.

Client/Server Architectures for DBMSs

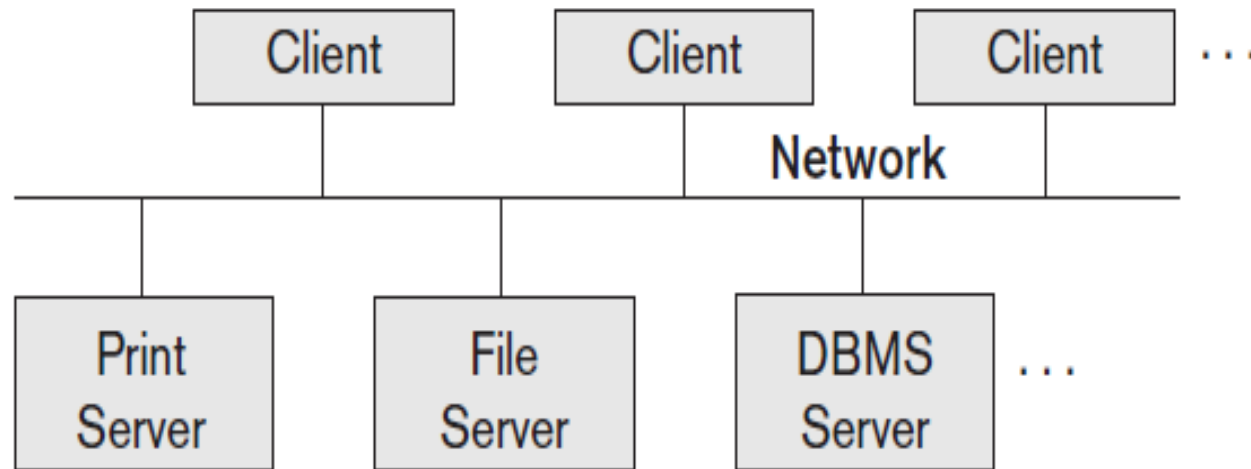


Figure 2.3 A Client/Server Architecture



Thank You