



# **Lecture 3**

## The Relational Data Model and Relational Database Constraints

# Agenda



Relational Model Concepts



Formal Definitions



Characteristics of Relations



Relational Model Constraints



Schema-based Constraints



Dealing with Constraint Violations

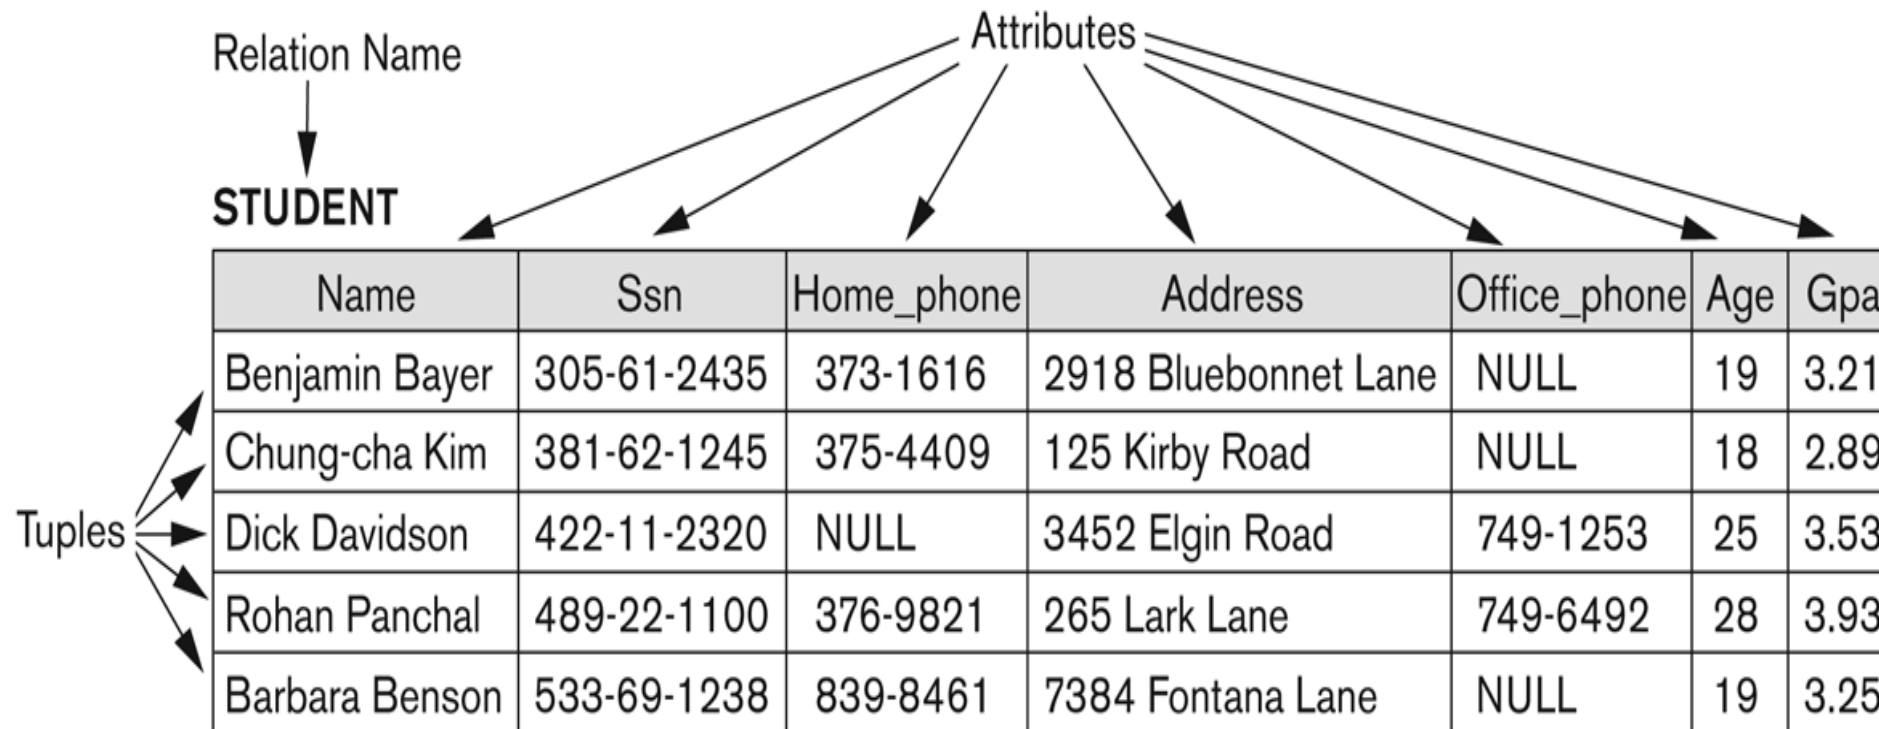


# Relational Model Concepts

# Relational Model Concepts

- The relational Model of Data is based on the concept of a *Relation*.
- A *Relation* looks like a **table** of values.
- A relation typically contains a **set of rows**.
- The data elements in each **row** represent certain facts that correspond to a real-world **entity** or **relationship**
  - Rows are called **tuples**
- Each **column** has a column header that gives an indication of the meaning of the data items in that column
  - The column header is called an **attribute name** (or just **attribute**)

# Relational Model Concepts



**Figure 5.1**

The attributes and tuples of a relation STUDENT.



# Formal Definitions

# Formal Definitions - Schema

➤ The **Schema** (or description) of a Relation:

- Denoted by **R(A1, A2, .....An)**
- R is the **name** of the relation
- The **attributes** of the relation are A1, A2, ..., An

➤ Example:

CUSTOMER (Cust-id, Cust-name, Address, Phone#)

- CUSTOMER is the relation name
- Defined over the four attributes: Cust-id, Cust-name, Address, Phone#

➤ Each attribute has a **domain** or a set of valid values.

- For example, the domain of Cust-id is 6 digit numbers.

# Formal Definitions - Tuple

- A **tuple** is an ordered set of values (enclosed in angled brackets ‘< ... >’)
- Each value is derived from an appropriate *domain*.
- Example:
  - A row in the CUSTOMER relation is a 4-tuple and would consist of four values:
    - <632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">
    - This is called a 4-tuple as it has 4 values
    - A tuple (row) in the CUSTOMER relation.
- A relation is a **set** of such tuples (rows)



# Formal Definitions - Domain

- A **domain** has a logical definition:
  - Example: “USA\_phone\_numbers” are the set of 10 digit phone numbers valid in the U.S.
- A **domain** also has a data-type or a format defined for it.
  - The USA\_phone\_numbers may have a format: **(ddd)ddd-dddd** where each d is a decimal digit.
  - Dates have various formats such as year, month, date formatted as **yyyy-mm-dd**, or as **dd mm,yyyy** etc.
- The attribute name designates the role played by a domain in a relation:
  - Used to interpret the meaning of the data elements corresponding to that attribute
  - Example: The domain Date may be used to define two attributes named “Invoice-date” and “Payment-date” with different meanings

## Formal Definitions - State

- A **Relation state**  $r$  of the relation schema  $R(A_1, A_2, \dots, A_n)$ , is a set of  $n$ -tuples  $r = \{t_1, t_2, \dots, t_m\}$ .
- Each  **$n$ -tuple**  $t$  is an ordered list of  $n$  values  $t = \langle v_1, v_2, \dots, v_n \rangle$ , where each value  $v_i$ ,  $1 \leq i \leq n$ , is an element of  $\text{dom}(A_i)$  or is a special NULL value.
- $n$  is the number of values within each tuple while  $m$  is the total number of tuples (or rows).

# Formal Definitions - Summary

<u>Informal Terms</u>	<u>Formal Terms</u>
Table	Relation
Column Header	Attribute
All possible Column Values	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	State of the Relation



# Characteristics of Relations

# Characteristics of Relations

- **Ordering of Tuples in a Relation.** A relation is defined as a *set* of tuples. Mathematically, elements of a set have *no order* among them; hence, tuples in a relation do not have any particular order. In other words, a relation is not sensitive to the ordering of tuples.
- **Ordering of Values within a Tuple.** According to the definition of a relation, an  $n$ -tuple is an *ordered list* of  $n$  values, so the ordering of values in a tuple is important.

# Characteristics of Relations

- **Values in the Tuples.** Each value in a tuple is an **atomic** value; that is, it is not divisible into components within the framework of the basic relational model. Hence, composite and multivalued attributes are not allowed.



# Relational Model Constraints

# Relational Model Constraints

➤ Constraints on databases can generally be divided into three main categories:

1. Constraints that are inherent in the data model. We call these **inherent model-based constraints** or **implicit constraints**.
2. Constraints that can be directly expressed in the schemas of the data model. We call these **schema-based constraints** or **explicit constraints**.
3. Constraints that *cannot* be directly expressed in the schemas of the data model, and hence must be expressed and enforced by the application programs or in some other way. We call these **application-based** or **semantic constraints** or **business rules**.



# Relational Model Constraints

## 1. **Inherent model-based constraints:**

- The **characteristics of relations** that we discussed are the inherent constraints of the relational model and belong to the first category.
- For example, the constraint that a relation cannot have duplicate tuples is an inherent constraint (this is because a relation is a *set* of tuples).

## 2. **Schema-based constraints:**

- The schema-based constraints include **Domain constraints**, **Key constraints**, **constraints on NULLs**, **Entity Integrity constraints**, and **Referential Integrity constraints**.

# Relational Model Constraints

## 3. **Semantic constraints:**

- Constraints in the third category are more general, relate to the **meaning** as well as **behavior** of attributes, and are difficult to express and enforce within the data model, so they are usually checked within the application programs that perform database updates.

Lets now discuss the **Schema-based constraints**



# Schema-based Constraints

## Schema-based Constraints - Domain Constraints

- Domain constraints specify that within each tuple, the value of each attribute  $A$  must be an atomic value from the domain  $\text{dom}(A)$ .
- The data types associated with domains typically include standard numeric data types for integers (such as short integer, integer, and long integer) and real numbers (float and double-precision float).
- Characters, Booleans, fixed-length strings, and variable-length strings are also available, as are date, time, timestamp, and other special data types.
- Domains can also be described by a subrange of values from a data type or as an enumerated data type in which all possible values are explicitly listed.

# Schema-based Constraints - Key Constraints

## ➤ Superkey of R:

- Is a set of attributes SK of R with the following condition:
  - No two tuples in any valid relation state  $r(R)$  will have the same value for SK
  - That is, for any distinct tuples  $t1$  and  $t2$  in  $r(R)$ ,  $t1[SK] \neq t2[SK]$
  - This condition must hold in *any valid state*  $r(R)$

## ➤ Key of R:

- A "minimal" superkey
- That is, a key is a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey (does not possess the superkey uniqueness property)

# Schema-based Constraints - Key Constraints

- Hence, a key is a superkey but not vice versa. A superkey may be a key (if it is minimal) or may not be a key (if it is not minimal).
- Consider the STUDENT relation of Figure 5.1. The attribute set {Ssn} is a key of STUDENT because no two student tuples can have the same value for Ssn.
- Any set of attributes that includes Ssn—for example, {Ssn, Name, Age}—is a superkey. However, the superkey {Ssn, Name, Age} is not a key of STUDENT because removing Name or Age or both from the set still leaves us with a superkey.

# Schema-based Constraints - Key Constraints

➤ In general:

- Any *key* is a *superkey* (but not vice versa)
- Any set of attributes that *includes a key* is a *superkey*
- A *minimal* superkey is also a key

➤ Hence, a key satisfies two properties:

1. Two distinct tuples in any state of the relation cannot have identical values for (all) the attributes in the key. This *uniqueness* property also applies to a superkey.
2. It is a *minimal superkey*—that is, a superkey from which we cannot remove any attributes and still have the uniqueness constraint hold. This *minimality* property is required for a key but is optional for a superkey.

# Schema-based Constraints - Key Constraints

- We call such key a **Primary Key**.
- If a relation has several **candidate keys**, one is chosen arbitrarily to be the *primary key*.
  - The primary key attributes are underlined.
- The *primary key* value is used to *uniquely identify* each tuple in a relation
  - Provides the tuple identity



# Schema-based Constraints - Key Constraints

➤ Example: Consider the CAR relation:

- It has two candidate keys: License\_number and Engine\_serial\_number.
- We choose License\_number (i.e. SerialNo) as the **primary key**

**CAR**

<u>License_number</u>	Engine_serial_number	Make	Model	Year
Texas ABC-739	A69352	Ford	Mustang	02
Florida TVP-347	B43696	Oldsmobile	Cutlass	05
New York MPO-22	X83554	Oldsmobile	Delta	01
California 432-TFY	C43742	Mercedes	190-D	99
California RSK-629	Y82935	Toyota	Camry	04
Texas RSK-629	U028365	Jaguar	XJS	04

**Figure 5.4**

The CAR relation, with two candidate keys: License\_number and Engine\_serial\_number.

# Schema-based Constraints - **NULLs**

- Another constraint on attributes specifies whether NULL values are or are not permitted.
- For example, if every STUDENT tuple must have a valid, non-NULL value for the Name attribute, then Name of STUDENT is constrained to be NOT NULL.

# Schema-based Constraints - Entity Integrity

- A **relational database schema**  $S$  is a set of relation schemas  $S = \{R_1, R_2, \dots, R_m\}$  and a set of **integrity constraints** IC.
- A **relational database state** must satisfy the integrity constraints specified in IC.
- Figure 5.5 shows a relational database schema that we call **COMPANY** = {EMPLOYEE, DEPARTMENT, DEPT\_LOCATIONS, PROJECT, WORKS\_ON, DEPENDENT}.
- In each relation schema, the underlined attribute represents the primary key.

# Schema-based Constraints - Entity Integrity

## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

## DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

## PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

## WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

## DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

**Figure 5.5**  
Schema diagram for  
the COMPANY  
relational database  
schema.

# Schema-based Constraints - Entity Integrity

## ➤ Entity Integrity:

- The *primary key attributes* PK of each relation schema R in S cannot have null values in any tuple.
  - This is because primary key values are used to *identify* the individual tuples.
  - $t[PK] \neq \text{null}$  for any tuple t
  - If PK has several attributes, null is not allowed in any of these attributes
- Note: Other attributes of R may be constrained to disallow null values, even though they are not members of the primary key.

## Schema-based Constraints - Referential Integrity

- A constraint involving **two** relations
  - The previous constraints involve a single relation.
- Used to specify a **relationship** among tuples in two relations:
  - The **referencing relation** and the **referenced relation**.

# Schema-based Constraints - Referential Integrity

- Tuples in the **referencing relation** R1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the **referenced relation** R2.
  - A tuple t1 in R1 is said to **reference** a tuple t2 in R2 if  $t1[FK] = t2[PK]$ .
- A referential integrity constraint can be displayed in a relational database schema as a directed arc from R1.FK to R2.

# Schema-based Constraints - Referential Integrity

## ➤ Referential Integrity:

- The value in the *foreign key* column (or columns) FK of the **referencing relation** R1 can be **either**:
  - (1) a value of an existing primary key value of a corresponding primary key PK in the **referenced relation** R2, or
  - (2) a **null**.

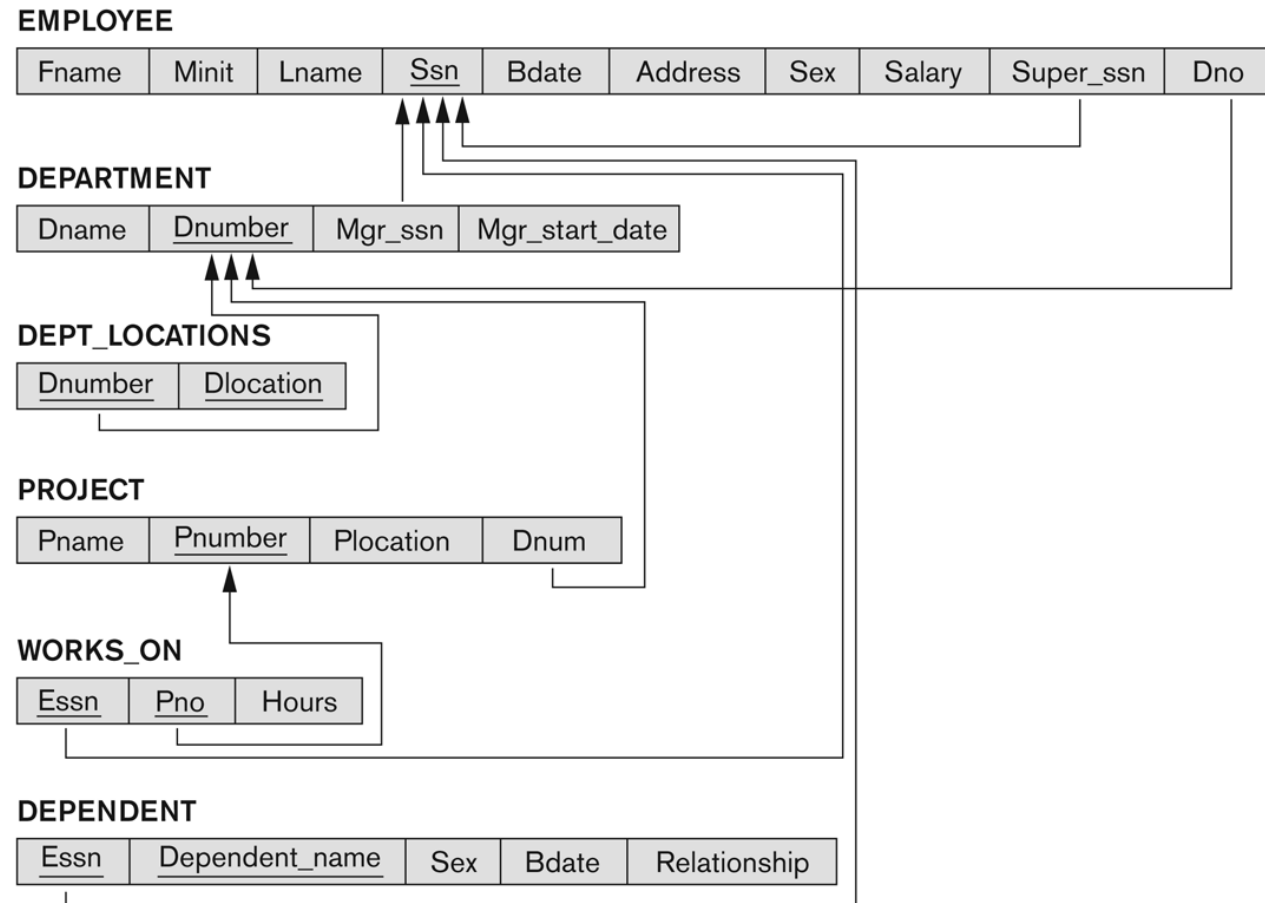
➤ In case (2), the FK in R1 should **not** be a part of its own primary key.



# Schema-based Constraints - Referential Integrity

**Figure 5.7**

Referential integrity constraints displayed on the COMPANY relational database schema.





# Dealing with Constraint Violations

# Constraint Violations

- Each *relation* will have many tuples in its current relation state
- Whenever the database is changed, a new state arises
- Basic operations for changing the database:
  - INSERT a new tuple in a relation
  - DELETE an existing tuple from a relation
  - UPDATE an attribute of an existing tuple
- Next slide shows an example state for the COMPANY database

Figure 5.6

One possible database state for the COMPANY relational database schema.

**EMPLOYEE**

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

**DEPARTMENT**

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

**DEPT\_LOCATIONS**

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

**WORKS\_ON**

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

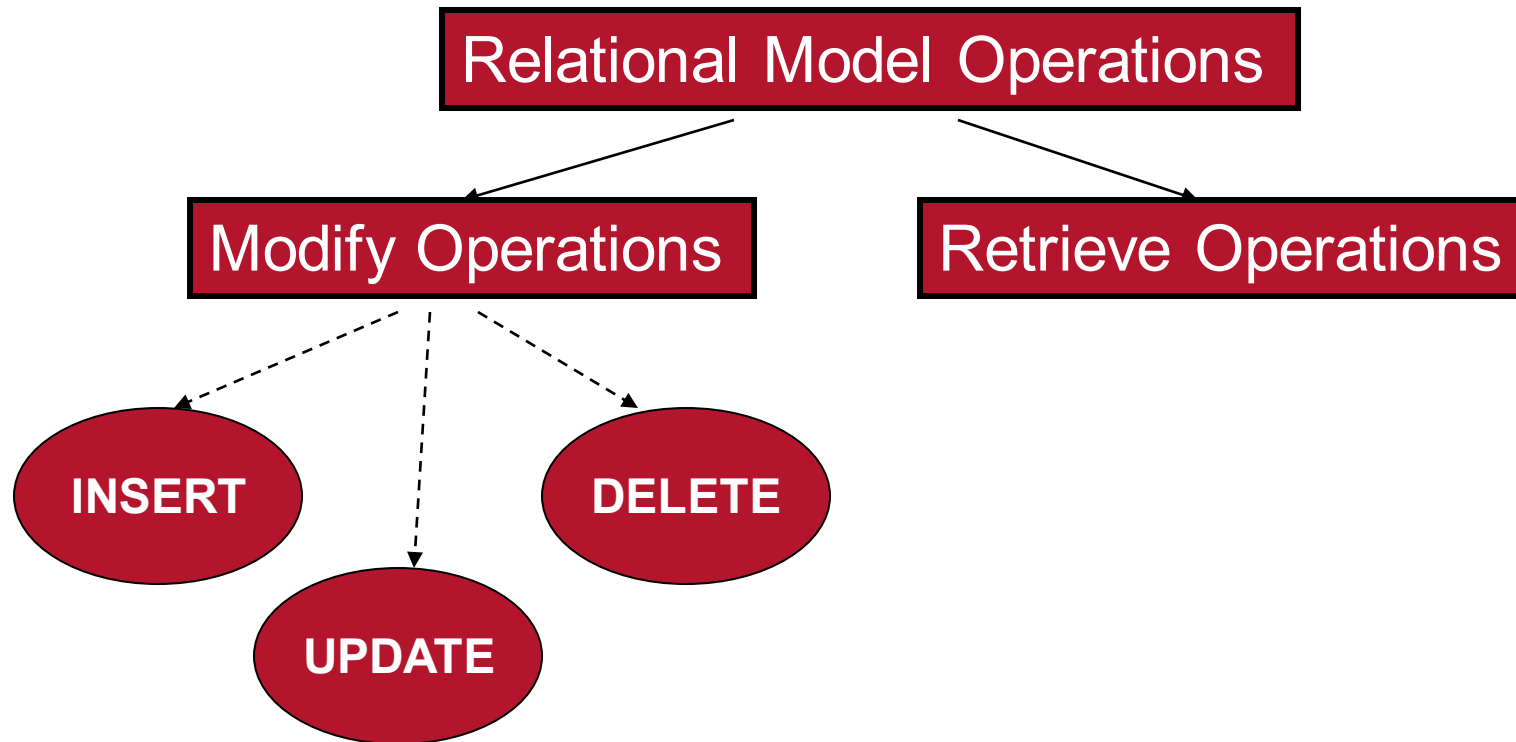
**PROJECT**

Pname	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

**DEPENDENT**

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

# Constraint Violations



Whenever modify operations are applied, the ICs should not be violated

# Constraint Violations

- In case of integrity violation, several actions can be taken:
- Cancel the operation that causes the violation (**RESTRICT** or **REJECT** option)
  - Perform the operation but inform the user of the violation
  - Trigger additional updates so the violation is corrected (**CASCADE** option, **SET NULL** option)
  - Execute a user-specified error-correction routine

# Possible violations for each operation

## ➤ INSERT may violate any of the constraints:

- Domain constraint:
  - if one of the attribute values provided for the new tuple is not of the specified attribute domain
- Key constraint:
  - if the value of a key attribute in the new tuple already exists in another tuple in the relation
- Referential integrity:
  - if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation
- Entity integrity:
  - if the primary key value is null in the new tuple

# Possible violations for each operation

## ➤ DELETE may violate only referential integrity:

- If the primary key value of the tuple being deleted is referenced from other tuples in the database
  - Can be resolved by several actions: RESTRICT, CASCADE, SET NULL
    - **RESTRICT** option: reject the deletion
    - **CASCADE** option: propagate the deletion by deleting tuples that reference the tuple that is being deleted.
    - **SET NULL** option: set the foreign keys of the referencing tuples to NULL
- One of the above options must be specified during database design for each foreign key constraint



# Possible violations for each operation

- UPDATE is similar to a DELETE operation followed by an INSERT operation.
- Therefore, UPDATE can violate any constraints violated by INSERT or DELETE in the previous slides.
- Constraints may be violated, depending on the attribute being updated:
  - Updating the primary key (PK):
    - May violate domain, key, referential integrity, entity integrity constraints
  - Updating a foreign key (FK):
    - May violate referential integrity
  - Updating an ordinary attribute (neither PK nor FK):
    - May violate domain constraints



Thank You